

The Architecture of Autonomy: A Strategic Analysis of AWS Agentic AI Services for Enterprise Business Transformation

The year 2026 marks a decisive pivot in the trajectory of industrial artificial intelligence, transitioning from the era of generative assistance to the era of agentic autonomy. While the preceding years were defined by Large Language Models (LLMs) that prioritized the generation and summarization of text, the current landscape is dominated by autonomous agents capable of reasoning, planning, and executing complex, multi-step tasks across disparate systems.¹ As enterprise value shifts from content generation to task accomplishment, Amazon Web Services (AWS) has engineered a vertically integrated ecosystem designed to support the scale, speed, and reliability demanded by these systems.¹ The transition is not merely technological but structural; industry forecasts suggest that by 2026, roughly 40% of all enterprise applications will include embedded task-specific AI agents, a significant increase from less than 5% in 2024.³

The strategic importance of agentic AI lies in its ability to function as a "true extension" of a human team, working autonomously for hours at a scale previously unattainable by manual labor or traditional automation scripts.² For the modern enterprise, the adoption of these agents represents an opportunity to solve the "context fragmentation" and "cognitive load" that currently plague knowledge workers, who are frequently forced to switch between dozens of interfaces and manually stitch together outputs.⁴ AWS has categorized its response to this demand into five distinct layers: Frontier Agents for team extension, ready-to-use Purpose-Built Applications, tools for building Custom Agents, cutting-edge Foundation Models, and a robust Infrastructure Foundation.²

I. Frontier Agents: Extending the Enterprise Team with Autonomous Logic

Frontier agents represent a new class of AI entities that deliver complete outcomes independently. Unlike simple chatbots that respond to a single prompt, these agents maintain persistent context across sessions and repositories, learning from feedback to handle increasingly complex assignments.¹ AWS has initially focused this capability on the core pillars of the digital enterprise: software development, security, and operations.²

The Kiro Autonomous Agent: Reimagining the Developer Workflow

In the software development lifecycle, the Kiro autonomous agent serves as an asynchronous partner that transforms how applications are built.² The primary capability of Kiro is its capacity to handle background development tasks, allowing human developers to focus on high-impact architectural decisions rather than routine feature implementation or bug fixing.¹ By maintaining a deep, persistent understanding of the codebase across multiple repositories, Kiro can translate high-level requirements into functional pull requests, including the necessary tests and documentation.¹

The business benefit of Kiro is a massive increase in development team capacity. Organizations using Kiro report an enhanced ability to say "yes" to a higher volume of product requests without a linear increase in headcount.² Furthermore, Kiro integrates with Model Context Protocol (MCP) servers, allowing it to gain instant expertise in specialized third-party tools such as Datadog, Figma, and Postman.⁶ This integration enables the agent to reason about the entire development ecosystem, from design handoffs to production monitoring.⁶

AWS Security and DevOps Agents: Proactive Governance and Resilience

The AWS Security Agent addresses the critical need for "security from the start" by embedding deep expertise into the design and deployment phases of an application.² This agent proactively reviews architectural design documents, scans every pull request against organizational and industry policies, and executes on-demand penetration testing to identify vulnerabilities before code reaches production.⁶ This shift-left approach reduces the cost of security remediation and minimizes the risk of compliance failures in highly regulated industries.

Parallel to this, the AWS DevOps Agent functions as an autonomous on-call engineer.⁶ Its core competency is the resolution and proactive prevention of system incidents.² By correlating telemetry data across observability platforms, source code, and CI/CD pipelines, the agent can diagnose root causes of anomalies and implement fixes autonomously.⁶ In an environment where downtime can cost thousands of dollars per minute, the DevOps Agent's ability to maintain system reliability without human intervention is a significant driver of operational efficiency and revenue protection.¹

Frontier Agent	Core Capability	Business Outcome	Potential Reach by 2026
Kiro	Asynchronous software development and	Increased dev capacity; faster feature delivery. ²	80% of dev tasks involve AI assistance. ⁷

	multi-repo understanding. ¹		
AWS Security Agent	Design review, PR scanning, and autonomous pentesting. ⁶	Reduced risk of vulnerabilities; embedded compliance. ²	Standard in Fortune 500 CI/CD pipelines. ³
AWS DevOps Agent	Autonomous incident resolution and telemetry correlation. ²	Proactive reliability; reduced MTTR. ²	60% of disruptions resolved without humans by 2031. ⁸

II. Purpose-Built Agentic Applications: Accelerating Value for Legacy and Business Systems

For enterprises seeking a faster time-to-market, AWS offers ready-to-use applications that eliminate the need for infrastructure management, allowing business leaders to focus on the value provided by agentic logic rather than the underlying compute.² These services target two primary enterprise challenges: modernizing legacy systems and consolidating fragmented business workflows.²

AWS Transform: The Agentic Modernization of Technical Debt

AWS Transform is a specialized agentic service designed to modernize complex workloads, including .NET, mainframe, and VMware estates.² It utilizes AI agents to analyze infrastructure, plan migration waves, and refactor code, effectively acting as an intelligent orchestrator for the entire transformation journey.⁹ The service has evolved from the code refactoring capabilities of Amazon Q Developer into a comprehensive migration suite that includes a collaborative web workspace for teams to manage projects outside of the standard AWS Console.⁹

####.NET and Mainframe Modernization Capabilities

The .NET specialized agent automates the transition from legacy .NET Framework (versions 4.5 through 4.8) to cross-platform .NET 8 on Linux.¹⁰ This modernization path is critical for reducing licensing costs and improving application performance.¹⁰ For mainframe workloads, AWS Transform employs agents to decompose monolithic COBOL applications into manageable Java-based microservices.¹¹ By extracting business rules and generating comprehensive documentation, the agent reduces reliance on scarce subject matter experts (SMEs) and shortens assessment timelines from several months to just a few days.¹¹

VMware to AWS Migration Acceleration

For VMware environments, the AWS Transform agent automates discovery and dependency mapping, creating intelligent transformation plans with AI-driven recommendations.¹³ A significant technical hurdle in these migrations is the conversion of on-premises networking into cloud-native AWS infrastructure; AWS Transform accelerates this network conversion by up to 80x.¹³ This automation allows organizations to exit data centers faster, eliminate technical debt, and reduce VMware licensing fees significantly.¹³

AWS Transform Workload	Performance / Productivity Metric	Time to Market Impact
.NET Framework	85% productivity boost; up to 15x performance improvement. ¹⁰	Modernization time reduced by 70-90%. ¹⁰
Mainframe (COBOL)	33% reduction in project timelines; 25% less analysis effort. ¹⁰	Assessment reduced from months to days. ¹²
VMware Estate	60-80% reduction in manual effort; 100% discovery accuracy. ¹⁰	Discovery and wave planning accelerated by 10x. ¹³

Amazon Quick Suite: The Unified Agentic Teammate

Amazon Quick represents a fundamental shift in the workplace experience, offering a unified, AI-powered workspace where research, business intelligence, and automation converge.⁴ It is designed to solve the problem of fragmented toolsets, where users lose time switching between CRM, collaboration, and analytics applications.⁴

Modules and Functional Capabilities

The Quick Suite is composed of five integrated components that collaborate to provide an agentic experience for non-technical users⁵:

- **Quick Sight:** The analytics layer that enables natural language data interaction without the need for SQL expertise.⁵
- **Quick Research:** An agent that performs deep analysis across internal documents, enterprise databases, and the web to generate cited reports and competitive intelligence.⁵
- **Quick Flows:** A "no-code" automation tool that turns repetitive tasks—such as report generation or follow-up emails—into agentic workflows in minutes.⁵
- **Quick Automate:** An orchestration engine for more complex, multi-agent, mission-critical business workflows.⁵

- **Quick Index:** A unified knowledge index that prepares all enterprise data, structured and unstructured, for fast retrieval by agents.⁵

The suite's strength lies in its ability to invoke third-party agents from providers like Box, Canva, and PagerDuty from a single interface.⁴ For example, an operations lead can use Quick to pull incident data from PagerDuty, generate a summary in Canva, and query relevant technical documents in Box without ever leaving the Quick workspace.⁴ This "single pane of glass" for agentic interaction is expected to be a major driver of productivity, as IDC predicts that 80% of workplace applications will embed AI copilots by 2026.⁷

III. Building and Deploying Custom Agents: Tailored Autonomy at Scale

For developers and enterprises that require bespoke agentic experiences, AWS provides a robust set of services and SDKs to build, deploy, and operate agents securely. This layer of the stack is designed for flexibility, allowing organizations to maintain full control over the logic, tools, and data connections while benefiting from fully managed infrastructure.²

Amazon Bedrock AgentCore: The Industrial-Scale Agent Platform

Amazon Bedrock AgentCore is a comprehensive platform for the lifecycle management of AI agents. It consists of fully managed services that work independently or together, allowing developers to choose between open-source flexibility and enterprise-grade reliability.²

Technical Architecture and Runtime Features

The AgentCore Runtime provides a secure, serverless environment for executing agents and their tools.¹⁶ It is designed for high-concurrency workloads, offering fast cold starts and session isolation, which ensures that every user interaction remains private and secure.¹⁶ A critical feature for long-running agents is the AgentCore Memory service, which supports both episodic memory (learning from past interactions) and long-term memory that persists across multiple sessions.¹⁶

AgentCore also introduces the "Gateway" component, which acts as a secure translator between agents and external systems.¹⁶ It converts standard APIs into Model Context Protocol (MCP)-compatible tools, allowing agents to interact with legacy systems or third-party services with minimal code changes.¹⁶ To ensure these agents operate within safe boundaries, the AgentCore Policy service allows developers to author fine-grained business rules using natural language or the Cedar open-source policy language.¹⁶

AgentCore Component	Functional Purpose	Developer Integration Benefit
Runtime	Serverless agent execution environment. ¹⁶	Any model/framework support; built-in security isolation. ¹⁶
Gateway	Connects agents to tools and APIs via MCP. ¹⁶	Secure, standardized access to thousands of tools. ¹⁶
Memory	Context persistence across sessions. ¹⁶	Agents that learn and adapt based on past outcomes. ¹⁷
Evaluations	Data-driven performance assessment. ¹⁶	Visualized agent quality metrics in CloudWatch. ¹⁷
Policy	Authoring of agent boundaries and rules. ¹⁷	Natural language control over agentic actions. ¹⁷
Browser	Cloud-based browser runtime for agents. ¹⁶	Safe, isolated web interaction and automation. ¹⁸

Strands Agents: A Model-Driven Approach to Rapid Development

Strands Agents is an open-source SDK (available in Python and TypeScript) that simplifies agent development by following a model-driven approach.² The core innovation of Strands is that it eliminates the need for complex hardcoded orchestration logic; instead, it leverages the reasoning capabilities of foundation models to handle planning and tool selection autonomously.¹⁹

For a coder, Strands is remarkably easy to integrate, requiring only a few lines of code to create a production-ready agent.² The SDK implements an "agentic loop" where the model iteratively reads context, plans an action, calls a tool, and incorporates the results until it reaches a goal.²⁰ Because the model itself decides when to use tools, developers do not need to write complex state machines or condition checks.¹⁹ Strands supports handoffs between agents and "swarm-style" cooperation, where multiple specialized agents work together on a single task.²⁰ In one documented instance, an organization using Strands reduced an investigation process from 30 minutes to 45 seconds, resulting in a \$5 million saving in operational costs.²¹

Amazon Nova Act: The Reliable Computer Use Service

One of the most significant challenges in agentic AI is "computer use"—the ability of an agent to interact with human-centric user interfaces.²² Amazon Nova Act is a purpose-built service for building and managing fleets of agents that automate browser-based UI workflows.² Unlike traditional automation tools like Selenium, which are brittle and break when a website's CSS structure changes, Nova Act achieves over 90% reliability by relying on visual and semantic understanding.¹⁷

The Perceive-Act-Repeat Architecture

Nova Act is powered by a custom Nova 2 Lite model that was trained end-to-end with the orchestrator and actuators using reinforcement learning in synthetic "web gyms".¹⁸ Its operation follows a repeating cycle:

1. **Perceive:** The agent takes a screenshot of the browser and parses the Document Object Model (DOM) to understand the current state.¹⁸
2. **Reason:** The Nova 2 Lite model analyzes the visual state against the user's instructions to produce an action plan (e.g., "click the Submit button").¹⁸
3. **Act:** The agent executes low-level commands like clicking at specific (x, y) coordinates or typing text via a secure, cloud-hosted browser.¹⁸
4. **Verify:** After the action, the agent takes a new screenshot to validate the outcome and determines the next step.¹⁸

This service is ideal for automating web QA testing, repetitive form filling across healthcare or travel portals, and competitive data extraction from sources that lack APIs.²⁴ PGA TOUR has utilized Nova Act to automate critical website testing across different geographies, maintaining high engagement with less manual effort.²⁵

IV. Foundation Models: The Intelligence Engine for Autonomous Action

Foundation models serve as the reasoning engines that drive agentic behavior. For an agent to be effective, its underlying model must excel at reasoning, planning, and structured tool use.² Amazon Bedrock provides access to a diverse portfolio of industry-leading models optimized for these complex workflows.²

Claude 4.5: The State of the Art in Agentic Performance

Anthropic's Claude 4.5 (available in Sonnet and Opus classes) has set a new benchmark for autonomous agents. Claude Sonnet 4.5 is particularly noted for its coding capabilities and computer use performance.²⁶ On the SWE-bench Verified benchmark, which measures a model's ability to resolve real GitHub issues, Sonnet 4.5 achieved a record 77.2%, outperforming competitors in production-ready code generation.²⁶

For agent developers, Claude 4.5 introduces features like "tool search," which allows an agent to dynamically load only the tools it needs from a library of hundreds, saving significant token costs and preventing schema confusion.²⁸ Its 200K token context window (with a 1M option) allows agents to process entire technical repositories or litigation records in a single session.²⁷

Benchmark / Capability	Claude Sonnet 4.5	Claude Opus 4.5	Qwen 3 MoE
SWE-bench (Coding)	77.2%	80.9%	Strong (optimized for agentic coding)
OSWorld (Computer Use)	61.4%	Improved navigation	Competitive browser use
MMMU (Visual Reasoning)	77.8%	80.7%	High (multimodal native)
Architecture	Dense	Dense	Mixture-of-Experts (MoE)

Qwen and Amazon Nova: Efficiency and Domain Depth

The Qwen 3 series from Alibaba, now available on Bedrock, utilizes a Mixture-of-Experts (MoE) architecture, activating only a subset of parameters for each request to provide high performance with efficient inference.²⁹ Qwen 3 introduces a "hybrid thinking" mode, allowing developers to switch between a step-by-step reasoning mode for complex problems and a rapid-response mode for simpler tasks.²⁹ This flexibility is essential for managing the cost-performance trade-offs in large-scale agent deployments.³⁰

The Amazon Nova family provides a tiered approach to intelligence. Nova Pro is designed for advanced reasoning, while Nova Premier acts as a "teacher" model used for distillation via Amazon Nova Forge.²² Forge allows organizations to create smaller, specialized models trained on proprietary data, ensuring that agents have deep, domain-specific intelligence without the latency of massive general-purpose models.²²

V. Infrastructure Foundation: Engineering for Scale and Reliability

The foundation layer is where AWS differentiates its agentic AI offering through custom silicon and integrated data centers. As agentic AI moves into production, the demand for

performance-per-dollar and low-latency inference becomes the dominant factor in business ROI.¹

Trainium and Inferentia: Custom Silicon for the Agentic Era

AWS Trainium3 and Inferentia2 chips are purpose-built for the massive scale of AI training and inference.² The Trainium3 UltraServer, launched in late 2025, matches the rack-scale performance of NVIDIA's Blackwell architecture (0.36 ExaFLOPS) while delivering up to 50% lower training costs.³¹ This silicon has been validated by Anthropic, which deployed over 500,000 Trainium chips to build "Project Rainier," the world's largest operational AI supercomputer for training the Claude model family.¹

For inference, the economic case is equally compelling. Customers report up to a 4x improvement in inference speed for generative video at half the cost when using AWS custom chips compared to traditional GPUs.³¹ This efficiency is critical for agents that perform thousands of "perceive-act" loops in the background.³¹

Amazon SageMaker AI and AI Factories

Amazon SageMaker AI provides the fully managed infrastructure to build, train, and deploy models at scale.² To support organizations with strict compliance or data sovereignty requirements, AWS offers "AI Factories"—dedicated AI infrastructure (NVIDIA or Trainium) deployed inside a customer's own data center.¹ This allows enterprises to build agentic systems that adhere to the most rigorous security standards while benefiting from the same performance as the public cloud.⁶

Infrastructure Component	Key Performance Value	Business Benefit
Trainium3 UltraServer	0.36 ExaFLOPS (Rack-scale parity with Blackwell). ³¹	50% lower training costs; 4x better energy efficiency. ³¹
Inferentia2	Optimized for real-time agent loops. ²	3x faster inference on Bedrock; 5x higher tokens per MW. ³¹
AI Factories	On-premises dedicated AI clusters. ⁶	Compliance and sovereignty for regulated industries. ⁶

Amazon S3 Vector	20 trillion vectors per bucket; <100ms query. ¹⁷	Scalable, real-time memory layer for global agent fleets. ¹⁷
------------------	---	---

VI. Business Value, Market Adoption, and Strategic Best Practices

The transition to agentic AI is no longer a future goal but a current operational standard. Market statistics for 2026 paint a picture of rapid adoption: the Agentic AI market is expected to reach \$10.86 billion in 2026, growing at a CAGR of 44.6% toward nearly \$100 billion by 2032.³

Popularity and Potential: The Surge of 2026

IDC expects AI copilots to be embedded in nearly 80% of enterprise workplace applications by 2026.⁷ Today, approximately 35% of organizations report broad usage of AI agents, with another 17% having rolled them out across the entire company.⁷ The potential for these systems is most evident in customer service, where Gartner predicts that agentic AI will autonomously resolve 80% of common issues by 2029, resulting in a 30% reduction in operational costs.³²

Business Benefits: Revenue and Productivity Gains

The primary driver for adoption is the tangible improvement in business metrics. In marketing, Gartner predicts that by 2028, 60% of brands will use agentic AI to provide "digital concierges" for hyper-personalized one-to-one interactions.⁸ In the software realm, the use of Kiro and AWS Transform has allowed companies like 1CloudHub and Landchecker to increase developer productivity by over 80% and reduce investigation times by 94%.¹⁰

The "Time to Market" for these solutions has shifted dramatically. What traditionally took months of development (such as mainframe modernization assessments or the creation of secure browser automation) can now be achieved in just days or weeks using AWS agentic services.¹²

Best Practices for Enterprise Integration

Analysis of successful AWS agentic deployments reveals several key best practices for business leaders:

- 1. Start with Spec-Driven Development:** Utilize Kiro and its specification-first approach to ensure that agents have a clear, validated roadmap before code is written, reducing rework and misaligned logic.²
- 2. Implement Cedar-Based Governance:** As agents gain more autonomy, utilize the Cedar policy language within AgentCore to define hard boundaries on what an agent can and

cannot do, ensuring security and compliance are maintained autonomously.¹⁷

3. **Leverage Multi-Agent Collaboration:** Don't rely on a single "god agent." Instead, use Strands or AgentCore to build a "swarm" of specialist agents that can call each other via A2A protocols, mimicking a high-performing human department.²⁰
4. **Prioritize Data Readiness:** Productivity loss of up to 15% is predicted for companies that fail to establish AI-ready data foundations.³ Use Amazon Quick Index to ensure that agents have high-speed, structured access to the organization's knowledge base.⁵

Conclusion: The Strategic Path to Autonomous Enterprise

The AWS Agentic AI ecosystem provides the first comprehensive, end-to-end framework for businesses to move beyond simple AI assistance into full operational autonomy. By splitting services into Frontier Agents, Purpose-Built Applications, Custom Development tools, Model Engines, and Infrastructure, AWS has created a roadmap that accommodates every level of technical maturity.²

For the enterprise, the message is clear: the breakthrough year for multi-agent systems has arrived.³³ The ability of agents like Nova Act to interact with UIs with 90% reliability, combined with the 50% cost savings offered by Trainium silicon, provides the economic and technical justification to scale these systems across every department.²³ Organizations that act now to integrate these agents as true extensions of their teams will not only see significant productivity gains but will also lead the next era of personalized, autonomous engagement in the global market.⁸ The future of business is not just powered by AI—it is executed by it.

Nguồn trích dẫn

1. AWS CEO calls AI inference a new building block that transforms what developers can build, truy cập vào tháng 3 25, 2026, <https://www.aboutamazon.com/news/aws/aws-ceo-ai-inference-transforms-developer-capabilities>
2. Agentic AI Solutions and Development Tools - AWS, truy cập vào tháng 3 25, 2026, <https://aws.amazon.com/ai/agentic-ai/>
3. Agentic AI Statistics 2026: Adoption, Market Size, Challenges & More - Cyntexa, truy cập vào tháng 3 25, 2026, <https://cyntexa.com/blog/agentic-ai-statistics/>
4. Amazon Quick adds third-party AI agents and expands built-in ..., truy cập vào tháng 3 25, 2026, <https://aws.amazon.com/about-aws/whats-new/2026/01/3p-agent-in-quick/>
5. Quick Suite — Meet Your New Agentic Teammate | by Rupal Bhatt ..., truy cập vào tháng 3 25, 2026, <https://engineering.doit.com/quick-suite-meet-your-new-agentic-teammate-dad0b32cd7fc>
6. 50+ New Announcements on re:Invent 2025 - DataOps Labs, truy cập vào tháng

3 25, 2026,

https://blog.dataopslabs.com/50-new-announcements-on-reinvent-2025?source=more_articles_bottom_blogs

7. The future of AI agents: Key trends to watch in 2026 - Salesmate CRM, truy cập vào tháng 3 25, 2026, <https://www.salesmate.io/blog/future-of-ai-agents/>
8. Gartner Predicts 60% of Brands Will Use Agentic AI to Deliver Streamlined One-to-One Interactions by 2028, truy cập vào tháng 3 25, 2026, <https://www.gartner.com/en/newsroom/press-releases/2026-01-15-gartner-predicts-60-percent-of-brands-will-use-agentic-ai-to-deliver-streamlined-one-to-one-interactions-by-2028>
9. A Cherry on Top of the AWS Migration Services: AWS Transform - Ippon Blog, truy cập vào tháng 3 25, 2026, <https://blog.ippon.tech/a-cherry-on-top-of-the-aws-migration-services-aws-transform>
10. Partners – AWS Transform – AWS, truy cập vào tháng 3 25, 2026, <https://aws.amazon.com/transform/partners/>
11. AWS Transform for mainframe, truy cập vào tháng 3 25, 2026, <https://aws.amazon.com/transform/mainframe/>
12. AWS for mainframe modernization – re:Invent 2025 Refresher ..., truy cập vào tháng 3 25, 2026, <https://aws.amazon.com/blogs/migration-and-modernization/aws-for-mainframe-modernization-reinvent-2025-refresher/>
13. Modernize VMware workloads with agentic AI - AWS Transform for VMware, truy cập vào tháng 3 25, 2026, <https://aws.amazon.com/transform/vmware/>
14. Your Agentic Teammate Enters the Chat: Amazon Quick Suite — Part 1 - Medium, truy cập vào tháng 3 25, 2026, <https://medium.com/@cloudelligent/your-agentic-teammate-enters-the-chat-amazon-quick-suite-part-1-118ee41681a8>
15. Activate Your Agentic AI Journey with Quick Suite - AWS Experience, truy cập vào tháng 3 25, 2026, <https://aws-experience.com/amer/smb/e/3cefa/activate-your-agentic-ai-journey-with-quick-suite>
16. Overview - Amazon Bedrock AgentCore - AWS Documentation, truy cập vào tháng 3 25, 2026, <https://docs.aws.amazon.com/bedrock-agentcore/latest/devguide/what-is-bedrock-agentcore.html>
17. AWS re:Invent 2025: AI Top Announcements | AWS Builder Center, truy cập vào tháng 3 25, 2026, <https://builder.aws.com/content/38WTIUxyJqj8vlgjcdHslQ64FAo/aws-reinvent-2025-ai-top-announcements>
18. Amazon Nova Act: How AWS Built a Reliable AI Browser Automation Agent (Architecture Deep Dive) | by Mohammed Anes - Medium, truy cập vào tháng 3 25, 2026, <https://medium.com/@mohammedanes008/amazon-nova-act-how-aws-built-a-reliable-ai-browser-automation-agent-architecture-deep-dive-39636a427e3c>

19. Beginner's Guide to Strands Agents: Getting Started with AI Agents on AWS, truy cập vào tháng 3 25, 2026, <https://builder.aws.com/content/375ld0kj8iqUxaR9rHRBMdf5ODX/beginners-guide-to-strands-agents-getting-started-with-ai-agents-on-aws>
20. Strands Agents SDK: A technical deep dive into agent architectures and observability - AWS, truy cập vào tháng 3 25, 2026, <https://aws.amazon.com/blogs/machine-learning/strands-agents-sdk-a-technical-deep-dive-into-agent-architectures-and-observability/>
21. Strands Agents — Open Source AI Agent SDK for Python & TypeScript, truy cập vào tháng 3 25, 2026, <https://strandsagents.com/>
22. Amazon Nova: Engineering the Future of Agentic AI - Tutorials Dojo, truy cập vào tháng 3 25, 2026, <https://tutorialsdojo.com/amazon-nova-engineering-the-future-of-agentic-ai/>
23. Build reliable AI agents for UI workflow automation with Amazon Nova Act, now generally available | AWS News Blog, truy cập vào tháng 3 25, 2026, <https://aws.amazon.com/blogs/aws/build-reliable-ai-agents-for-ui-workflow-automation-with-amazon-nova-act-now-generally-available/>
24. Amazon Nova Act, truy cập vào tháng 3 25, 2026, <https://nova.amazon.com/act>
25. Build reliable agents to automate production UI workflows at scale – Amazon Nova Act, truy cập vào tháng 3 25, 2026, <https://aws.amazon.com/nova/act/>
26. Introducing Claude Sonnet 4.5 - Anthropic, truy cập vào tháng 3 25, 2026, <https://www.anthropic.com/news/claude-sonnet-4-5>
27. Claude Sonnet 4.5: Features, Benchmarks & Pricing (2026) - Leanware, truy cập vào tháng 3 25, 2026, <https://www.leanware.co/insights/claude-sonnet-4-5-overview>
28. Claude Opus 4.5 now in Amazon Bedrock | Artificial Intelligence - AWS, truy cập vào tháng 3 25, 2026, <https://aws.amazon.com/blogs/machine-learning/claude-opus-4-5-now-in-amazon-bedrock/>
29. Qwen models are now available in Amazon Bedrock | AWS News Blog, truy cập vào tháng 3 25, 2026, <https://aws.amazon.com/blogs/aws/qwen-models-are-now-available-in-amazon-bedrock/>
30. Qwen - Models in Amazon Bedrock - AWS, truy cập vào tháng 3 25, 2026, <https://aws.amazon.com/bedrock/qwen/>
31. AWS Trainium3 Matches NVIDIA: 50% Cheaper AI Training | byteiota, truy cập vào tháng 3 25, 2026, <https://byteiota.com/aws-trainium3-matches-nvidia-50-cheaper-ai-training/>
32. Gartner Predicts Agentic AI Will Autonomously Resolve 80% of Common Customer Service Issues Without Human Intervention by 2029, truy cập vào tháng 3 25, 2026, <https://www.gartner.com/en/newsroom/press-releases/2025-03-05-gartner-predicts-agentic-ai-will-autonomously-resolve-80-percent-of-common-customer-service-issues-without-human-intervention-by-20290>
33. AI Agent Adoption 2026: What the Data Shows | Gartner, IDC - Joget, truy cập

vào tháng 3 25, 2026,

<https://joget.com/ai-agent-adoption-in-2026-what-the-analysts-data-shows/>